


In-place CartesianTreeSort variant

Francisco Alejandro Arganis Ramírez

NOTE. The following conventions are used:


- Array indices start at 1

The valid range of an index i for an array with size n is $1 \leq i \leq n$.



- Node heights start at 1

The height of a node is the number of nodes in the longest downward path to a leaf from that node.



The height of a tree its the height of its root.

CartesianTreeSort is an adaptive sorting algorithm with $O(n \log n)$ time and $O(n)$ space worst case complexities. The standard algorithm uses a min cartesian tree, which is an implicit binary tree placed in an array with the two additional properties

1. The tree is also a min heap.
2. An in-order traversal of the nodes yields the values in the same order in which they appear in the array.

CartesianTreeSort works by first building the min cartesian tree and then successively finding the next smallest value. This idea can be adapted to be done in-place, but the improvement in memory usage results in an non-adaptive algorithm.


The variant, refered to as **InplaceCartesianTreeSort**, uses a modified min cartesian tree, defined as follows: A min left oriented cartesian tree (loct) is an implicit binary tree placed in an array with the two additional properties

1. The tree is also a min heap.
2. An pre-order traversal of the nodes yields the values in the same order in which they appear in the array.

Alternatively, a right oriented cartesian tree (roct) that satisfies the post-order property could be defined instead. In this case a max roct is required to sort the array.

InplaceCartesianTreeSort replicates the behavior of **CartesianTreeSort**, by first building a min loct and then successively finding the next smallest value in total $O(n \log n)$ time and $O(1)$ space worst case. However, for almost-sorted arrays **CartesianTreeSort** will run in $O(n)$ time while **InplaceCartesianTreeSort** will still run in $O(n \log n)$.

The advantage of using pre-order over in-order is that node indices can always be directly calculated through a simple formula. The tree structure scheme for a loct and the corresponding index distribution on the implicit tree are shown below.



With this scheme a node at index i and height h has their left child at index $i + 1$ and right child at index $i + 2^{h-1}$.

leftChild. Receives an index i and returns the index of its left child.

```
leftChild ( i )
begin
    return i + 1
end
```

rightChild. Receives an index i with height h and returns the index of its right child.

```
rightChild ( i , h )
begin
    return i + (1 << (h - 1))
end
```

The min loc building requieres only $O(n)$ time, just like building a regular min cartesian tree, and it works in the same way as the regular **heapify** algorithm from **HeapSort**. First, the height H of the implicit tree is calculated.

`getMaxHeight`. Receives the size n of an array and returns the height of the implicit tree, that is, returns an integer H such that $2^{H-1} - 1 < n \leq 2^H - 1$.

```
getMaxHeight (n)
```

```

begin
    height <- floor (log2 (n))
    if (1 << height) = n + 1 do
        return height
    else do
        return height + 1
    end if
end
```

All nodes with height 1 are ignored as they are a valid loc by themselves. Then, for each height from 2 to H , the values at nodes with that height are sifted down.

`tripleMin`. Receives an array a and three indices i, j, k and returns an index m such that $a[m] = \min\{a[i], a[j], a[k]\}$.

```
tripleMin (a, i, j, k)
```

```

begin
    min <- i
    if j <= a.size && a [j] < a [min] do
        min <- j
    end if
    if k <= a.size && a [k] < a [min] do
        min <- k
    end if
    return min
end
```


`minSiftDown`. Receives an array a and an index i with height h and sifts down $a[i]$ to its appropriate position in the min loc.

```
minSiftDown (a, i, h)
```

```

begin
    while h > 1 do
        left <- leftChild (i)
        right <- rightChild (i, h)
        min <- tripleMin (a, i, left, right)
        if min = i do
            break
        end if
        a.swap (i, min)
        i <- min
        h <- h - 1
    end while
end
```

The difference of two consecutive nodes with the same height h is $r + 2^h$ where r is the number of roots between them with greater height. This can be calculated as $2^h - 1 + g$ where g is a gap given by the binary carry sequence (OEIS A007814). The first few values of g are 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, 4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, 5, ... and have a closed-form expression.



`binaryCarrySequence`. Receives an integer c and returns the exponent of highest power of 2 dividing c .

binaryCarrySequence (c)

```

begin
    return log2 (c - (c & (c - 1)))
end

```

`buildMinLoct`. Receives an array a and the height of the implicit tree H and builds a min loct.

buildMinLoct (a, H)

```


begin
    height <- 2
    while height <= H do
        counter <- 1
        increment <- (1 << height) - 1
        node <- 1 + H - height
        while node < a.size do
            minSiftDown (a, node, height)
            node <- node + increment + binaryCarrySequence (counter)
            counter <- counter + 1
        end while
        height <- height + 1
    end while
end

```

The chosen tree structure scheme converts a linear traversal of the region of any subtree in the array into a pre-order traversal of that subtree. Once the min loct is built, for each node visited during the traversal, the algorithm `findNextMin` compares the value at that index against the values at the roots of the largest possible right subtrees that appear later in the array. As there cannot be more than $\log n$ of these right subtrees, the operation takes $O(\log n)$.

If `findNextMin` finds the next minimum as the root of a right subtree, the minimum is extracted from its respective min loc and the original value is sifted down in $O(\log n)$. So the total time complexity of the traversal step is $O(n \log n)$.

Given an index i , the roots of the largest possible right subtrees after i are found by performing a binary-like search of i . Starting from the root, move to the right child j . If $j > i$ compare $a[i]$ against $a[j]$ and move to the left child. If $j < i$ then move to the right child. Repeat until $j = i$. This search keeps track of the height of the roots so that the sifting can be performed without additional calculations.



`findNextMin`. Receives an array a , an index i and the height of the implicit tree H and returns a pair (m, h) where m is an index of height h such that $a[m] \leq a[k]$ for all $k \geq i$.

findNextMin (a , i , H)

```

begin
    height <- H
    current_node <- 1
    m <- i
    h <- 0
    while current_node != i do
        next_node <- rightChild ( current_node , height )
        if next_node > i do
            if next_node <= a.size && array [ next_node ] < array [ m ] do
                m <- next_node
                h <- height - 1
            end if
            current_node <- leftChild ( current_node )
        else do
            current_node <- next_node
        end if
        height <- height - 1
    end while
    return (m, h)
end

```

`minLoctTraversal`. Receives an array a and the height of the implicit tree H and sorts it if the min loct has been built.

```
minLoctTraversal (a, H)
```

```
begin
    index <- 2
    while index < a.size
        (m, h) <- findNextMin (a, index, H)
        if m != index
            a.swap (index, m)
            minSiftDown (a, m, h)
        end if
        index <- index + 1
    end while
end
```

The pseudo-code for the complete algorithm is as follows.

`InplaceCartesianTreeSort`. Receives an array a and sorts it.

```
InplaceCartesianTreeSort (a)
```

```
begin
    H <- getMaxHeight (a.size)
    buildMinLoct (a, H)
    minLoctTraversal (a, H)
end
```

All the procedures `InplaceCartesianTreeSort` uses are iterative with constant memory requirements. As a result, the total space complexity is $O(1)$ worst case, just like `HeapSort`.